

Calico Ingress Gateway

Month Year

Presented by *Solutions Architect*



Agenda



- Recapitulative
- Configuration options
 - Traffic routing
 - Traffic shaping
 - Other traffic functions

01

Recapitulative

Kubernetes Native vs Ingress Gateway

- Kubernetes provides three core service types for exposing apps:
 - ClusterIP - Internally only
 - NodePort - Inflexible, manual, conflictual
 - LoadBalancer - Costly (1 per service)
- Native Kubernetes services handle L4 exposure but lack L7 flexibility, security, and portability. Ingress/Gateway API fills these gaps with standardised, feature-rich traffic management.
- Each solves a piece of the puzzle, but gaps remain.

Feature	Native Kubernetes (ClusterIP/NodePort/LoadBalancer)	Ingress/Gateway API
External Access	Limited (NodePort: manual port management, LoadBalancer: cloud-specific).	Standardized L7 (HTTP/HTTPS) routing.
Traffic Routing	L4 only (no host/path rules).	Path-based (/api), host-based (app.example.com), headers, traffic splitting.
TLS Termination	Manual cert management (e.g., Service Mesh needed).	Built-in TLS termination (e.g., cert-manager integration).
API Gateway Features	None (requires third-party tools).	Rate limiting, auth (OIDC/JWT), request rewriting, retries.
Multi-Tenancy	No namespace isolation for routing.	HTTPRoute scoped to namespaces (Gateway API).
Cloud Portability	Tied to cloud LoadBalancers (AWS ALB, GCP LB).	Vendor-neutral configuration.
Security	Basic NetworkPolicy (L3/L4).	mTLS, WAF integration, L7 policies.
Observability	Limited (service metrics only).	Envoy access logs, Prometheus metrics, distributed tracing.

02

Configurations options

Critical Gaps in Kubernetes Traffic Control

"How do you roll out new features without downtime or risk?"

- What happens if the new version has a critical bug?
- How can we minimize user impact during deployments?
- Is there a way to automatically roll back if errors occur?

"How do we handle failures gracefully without manual intervention?"

- What if a backend becomes unhealthy—does traffic automatically shift away?
- Can we define SLA-based routing (e.g., route away if latency > 500ms)?
- Is there circuit-breaking to prevent cascading failures?



"How do we support multi-cluster and hybrid-cloud routing?"

- Can we split traffic across clusters (e.g., 80% in AWS, 20% in GCP)?
- How do we handle failover if a cluster goes down?
- Is there geo-aware routing to direct users to the nearest region?

"How do we monitor and troubleshoot traffic routing in real-time?"

- Can we see live traffic distribution between versions?
- Are there detailed metrics (success rate, latency, errors) per route?
- How do we trace a request path for debugging?

Where Legacy Routing Falls Short



- **No native Kubernetes support for canary/blue-green deployments**
 - Manual scripting or external tools required.
 - No fine-grained control over traffic shifting (e.g., 5% → 10% → 50%).
- **Hard-coded routing rules in legacy Ingress**
 - Static configurations require redeployment for changes.
 - No dynamic traffic splitting based on conditions.
- **Lack of observability in traffic routing**
 - Difficulty tracking which requests go to which version.
 - No built-in metrics for success/failure rates per backend.
- **Security and compliance risks**
 - No built-in way to enforce zero-trust policies during traffic shifts.
 - Risk of exposing unfinished features to unintended users.

2.1

Traffic Routing

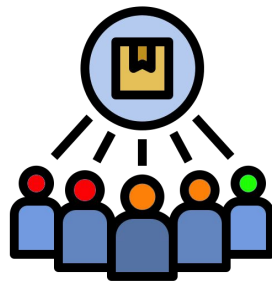
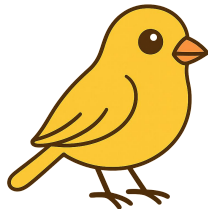
Solution: Traffic Splitting




1. **Weighted** Splitting
 - Percentage-based distribution (e.g., 70/30)
2. **Header-Based** Routing
 - Split based on HTTP headers (e.g., `x-user-type: premium`)
3. **Path-Based** Routing
 - Route by URL path (e.g., `/v1/` → `v1-service`)
4. **Query Parameter** Routing
 - Split by URL query params (e.g., `?version=beta`)
5. **Cookie-Based** Routing
 - Direct traffic using session cookies
6. **Geolocation** Routing
 - Route by client IP/country (e.g., EU → `europe-service`)
7. Traffic **Mirroring (Shadowing)**
 - Copy traffic to another cluster without affecting responses
8. **Canary** Deployments
 - Subset-based splitting (e.g., 5% to new version)
9. **Runtime Fractional** Routing
 - Dynamic splits controlled by runtime configuration
10. **Priority-Based** Routing
 - Route to clusters based on priority levels
11. **Load Balancer Subsets**
 - Split traffic to endpoint subsets (e.g., by zone/version)

Each type can be combined (e.g., **weighted + header-based****) for advanced use cases.**

Common Deployment Strategies with Calico Gateway




Canary Deployments



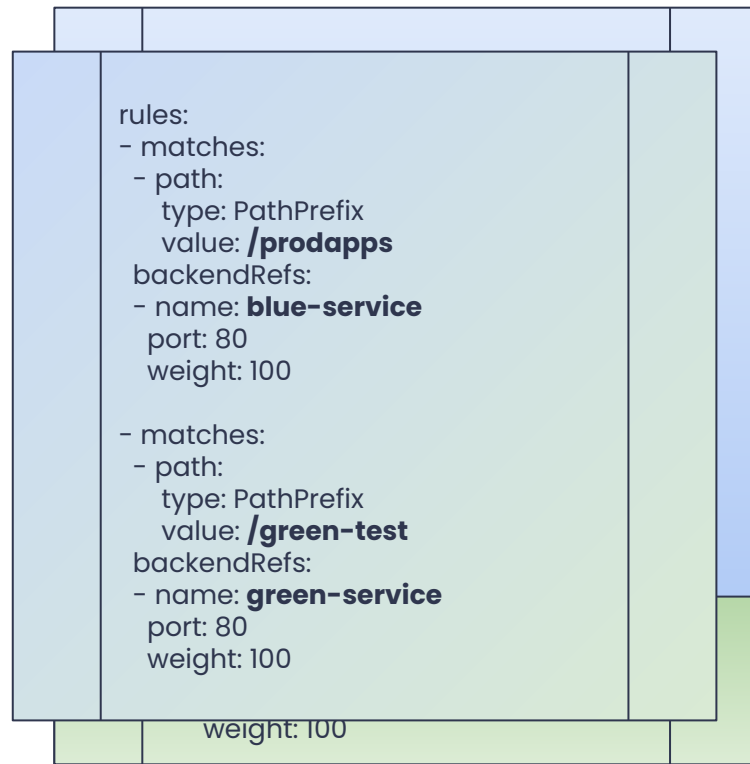
- Gradually roll out new versions to a subset of users
- Traffic splitting based on weight 90/10
- Often used for testing new features with minimal risk

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: canary-demo
  namespace: prodapps
spec:
  parentRefs:
    - name: calico-gateway
      kind: Gateway
  hostnames:
    - "prodapps.com"
  rules:
    - matches:
        - path:
            type: PathPrefix
            value: /prodapps
      backendRefs:
        - name: canary-service
          port: 80
          weight: 10
        - name: primary-service
          port: 80
          weight: 90
```

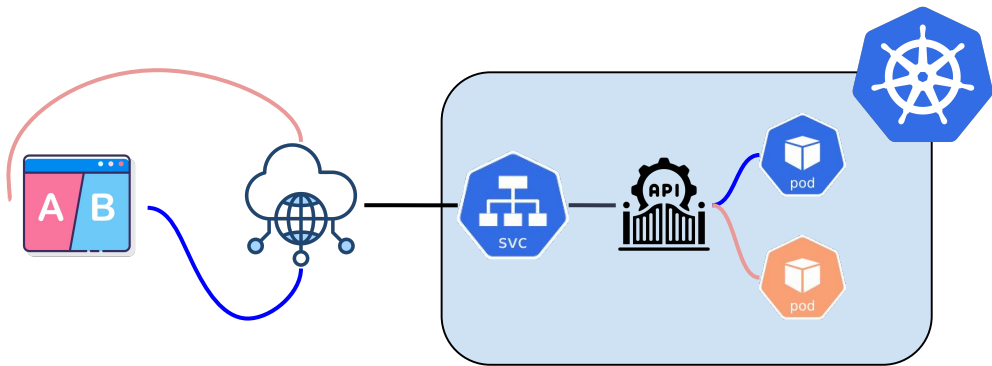
Blue Green Deployments



- Maintain two identical production environments
- Switch all traffic from "blue" (old) to "green" (new) at once
- Enables instant rollback if issues occur



A/B Testing

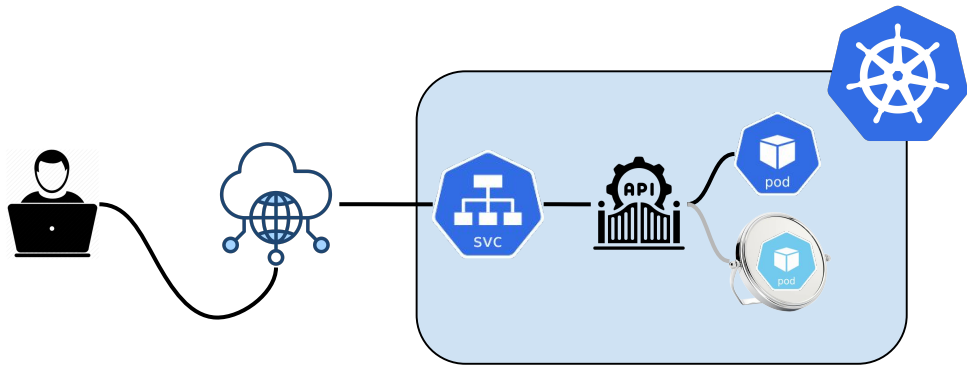


- Route traffic based on request attributes (headers, cookies, etc.)
- Useful for testing different versions with specific user segments (example: Mobile vs PC browsers)
- More sophisticated than simple percentage-based split

```
rules:
#B group: if header "x-ab-group: B" is
present
- matches:
- path:
  type: PathPrefix
  value: /prodapps
  headers:
  - name: x-ab-group
    value: B
  backendRefs:
  - name: ab-version-b
    port: 80
    weight: 100

#Default group A — all other traffic
- matches:
- path:
  type: PathPrefix
  value: /prodapps
  backendRefs:
  - name: ab-version-a
    port: 80
    weight: 100
```


Traffic Mirroring (Shadowing)



- Send a copy of production traffic to new version
- Doesn't affect live users while testing real traffic patterns
- Helps validate performance before actual cutover

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: mirroring-example
  namespace: default
spec:
  parentRefs:
    - name: calico-envoy-gateway
      kind: Gateway
      namespace: tigera-gateway
  hostnames:
    - "app.example.com"
  rules:
    - matches:
        - path:
            type: PathPrefix
            value: /myapp
      filters:
        - type: RequestMirror
          requestMirror:
            backendRef:
              name: shadow-service
              port: 8080
              group: ""
              kind: Service
            backendRefs:
              - name: primary-service
                port: 8080
```

Header-Based Routing



- Route traffic based on HTTP headers
- Useful for internal testing or feature flags
- Example: Route employees to new version while customers stay on stable

Rule 1: **Green users** (new features)

- matches:

- headers:

- type: **Exact**

name: **x-user-color**

value: **green**

filters:

- type: *RequestHeaderModifier*

requestHeaderModifier:

set:

- name: *x-routed-color*

value: *green*

backendRefs:

- name: **green-service**

port: 8080

weight: 100

Rule 2: **Yellow users** (beta features)

- matches:

- headers:

- type: **Exact**

name: **x-user-color**

value: **yellow**

backendRefs:

- name: **yellow-service**

port: 8080

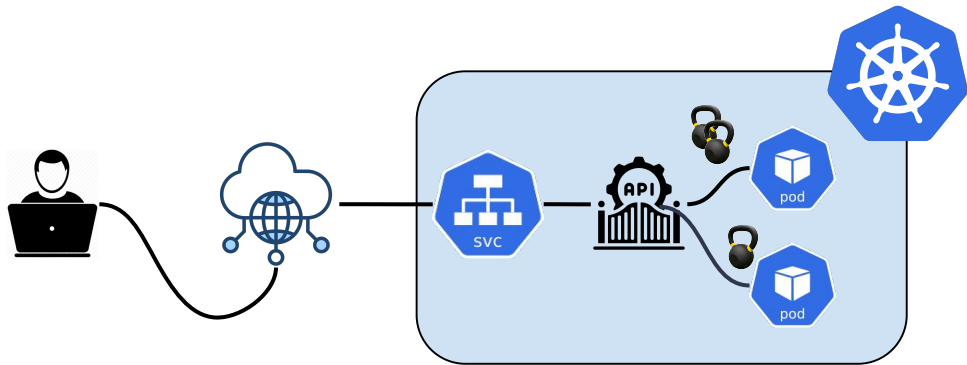
Rule 3: **Red users** (default stable)

- backendRefs:

- name: **red-service**

port: 8080

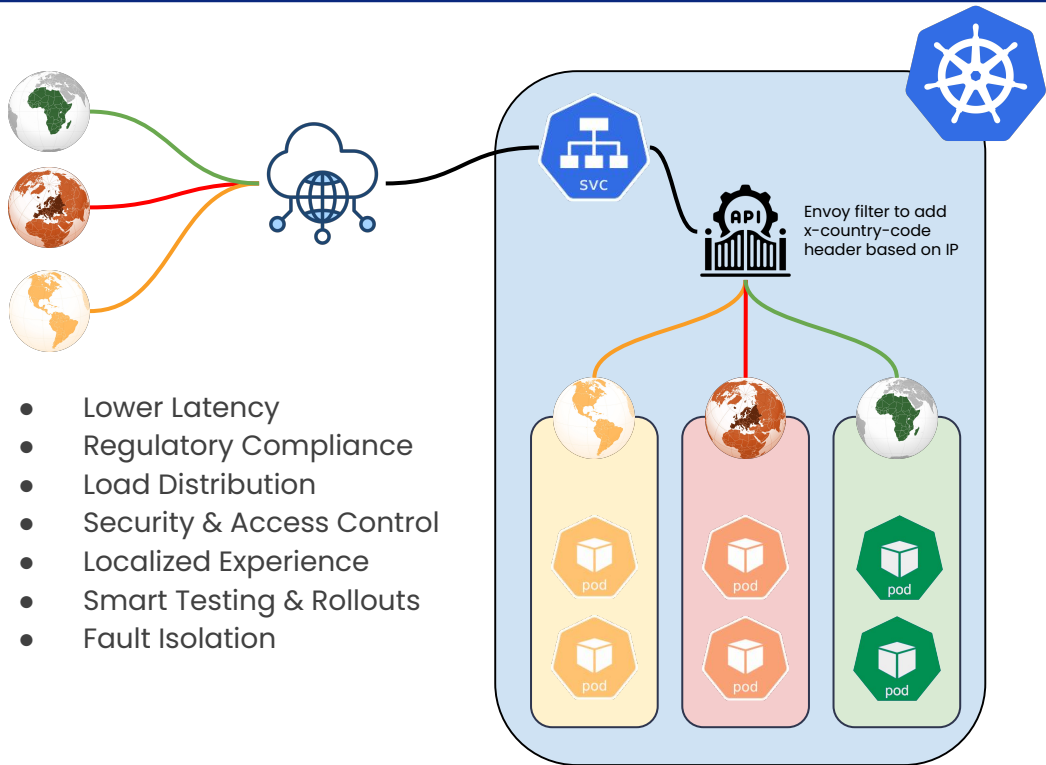
Weighted Traffic Splitting



- Distribute traffic across multiple service versions
- More flexible than simple canary (can split across multiple versions)
- Example: 70% v1 and 30% v2

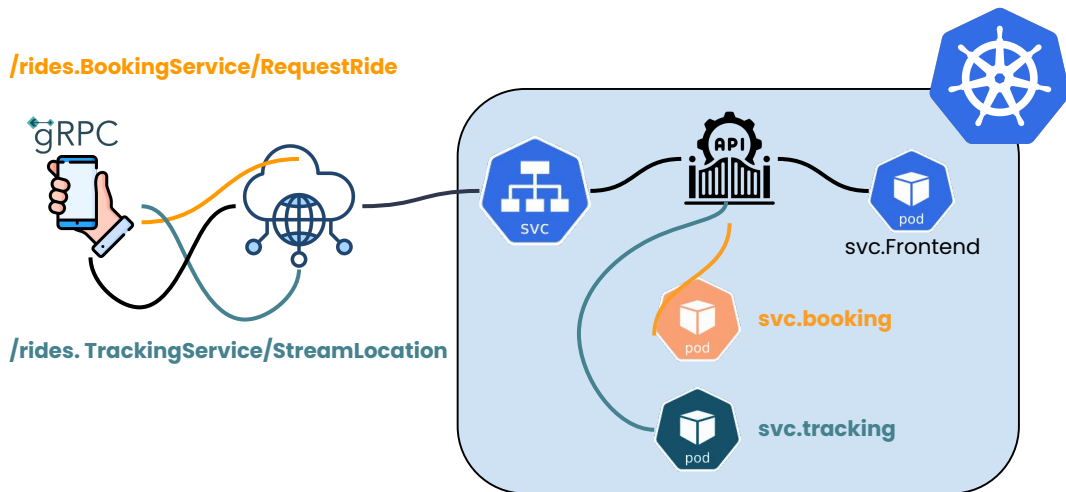
```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: basic-weighted-split
spec:
  parentRefs:
    - name: envoy-gateway
  rules:
    - backendRefs:
        - name: service-v1
          port: 80
          weight: 70
        - name: service-v2
          port: 80
          weight: 30
```


Geographic Routing



```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: geo-routing
spec:
  parentRefs:
  - name: envoy-gateway
  rules:
  - matches:
    - headers:
      - name: "x-geo-region"
        value: "africa"
    backendRefs:
      - name: africa-service
        port: 80
  - matches:
    - headers:
      - name: "x-geo-region"
        value: "europe"
    backendRefs:
      - name: europe-service
        port: 80
  - matches:
    - headers:
      - name: "x-geo-region"
        value: "america"
    backendRefs:
      - name: america-service
        port: 80
  - backendRefs: # Default route
    - name: default-service
      port: 80
```

gRPC Routing



- Your mobile app sends gRPC requests directly to Envoy Gateway.
- Envoy routes each method to the right microservice (no frontend involved).
- Result: Faster, more efficient, and scalable than traditional REST-through-frontend approaches.
- gRPC also supports all the previous http deployments we have covered and many more.

```
hostnames:  
- "api.rideshare.com"  
rules:  
- matches:  
  - method:  
    service: "rides.BookingService"  
    method: "RequestRide"  
  backendRefs:  
    - name: booking-service  
      port: 50051  
      weight: 100 # 100% to  
        booking-service  
  
- matches:  
  - method:  
    service: "rides.TrackingService"  
    method: "StreamLocation"  
  backendRefs:  
    - name: tracking-service  
      port: 50052  
      weight: 100 # 100% to  
        tracking-service
```

2.2

Traffic Shaping

Traffic Shaping

Rate Limiting

Global (external service) and local (in-process) request throttling.
Supports requests per second (RPS) or connection limits.

Circuit Breaking

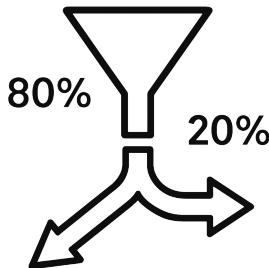
Automatically blocks traffic to overwhelmed services (based on errors/timeouts).
Configurable thresholds for max connections, pending requests, etc.

Retry Policies

Controls retry attempts for failed requests (with backoff strategies).
Can filter retries based on status codes/gRPC codes.

Timeouts

Sets deadlines for requests (global, per-route, or per-cluster).



Load Shedding

Drops or queues requests when upstream services are overloaded.

Request Buffering

Delays or buffers requests (e.g., for streaming or batch processing).

Bandwidth Limits

Throttles bandwidth for HTTP/TCP streams (bytes per second).

Fault Injection

Simulates failures (aborts/delays) to test resilience.

Priority-Based Routing

Directs traffic to high/low-priority clusters based on load.

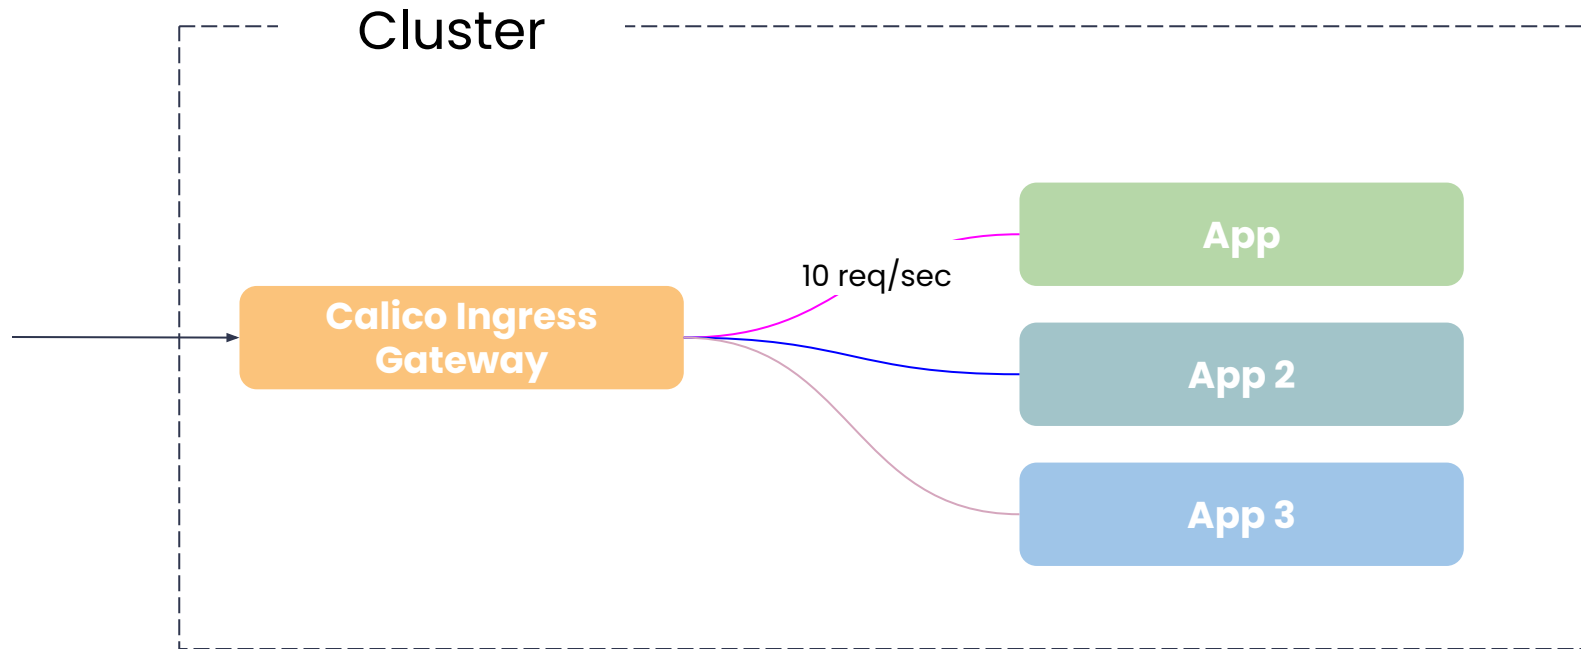
Adaptive Concurrency

Dynamically adjusts request limits based on latency metrics.

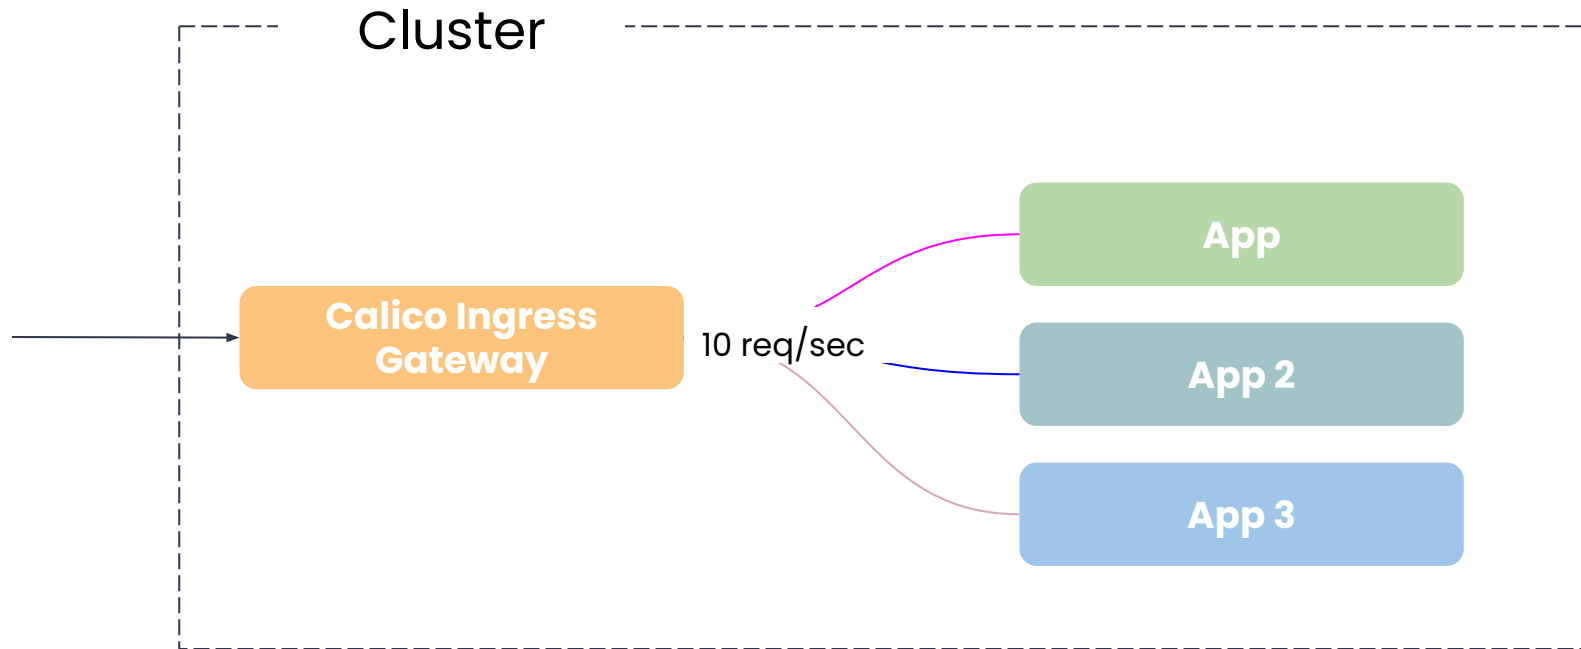
Quota Management

Enforces usage quotas (e.g., API calls/user).

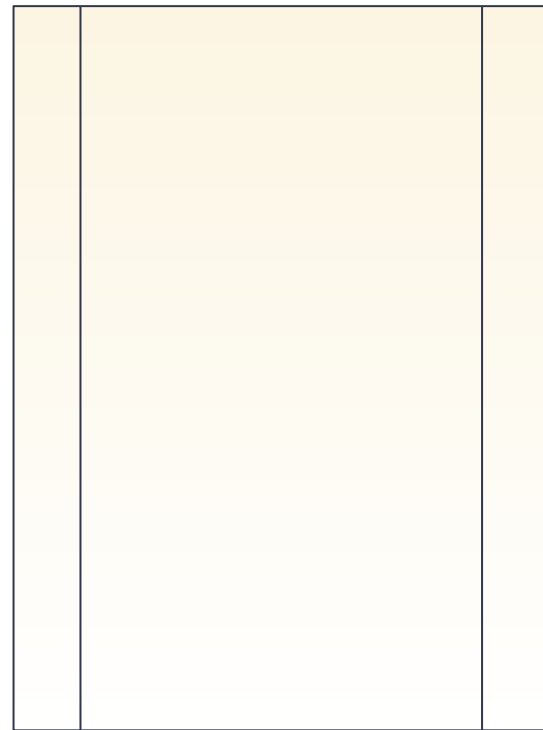
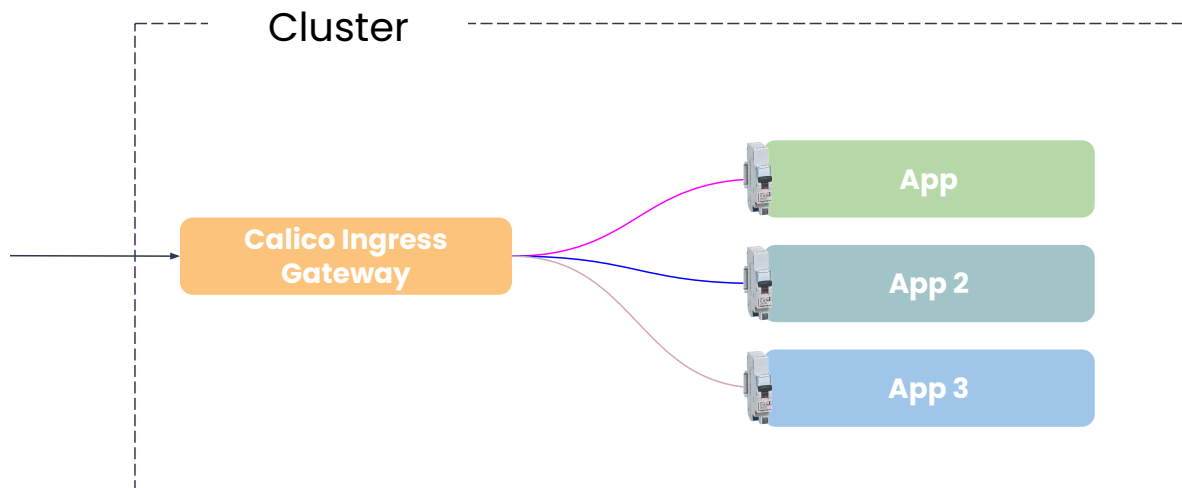
Local Rate Limiting



Global Rate Limiting



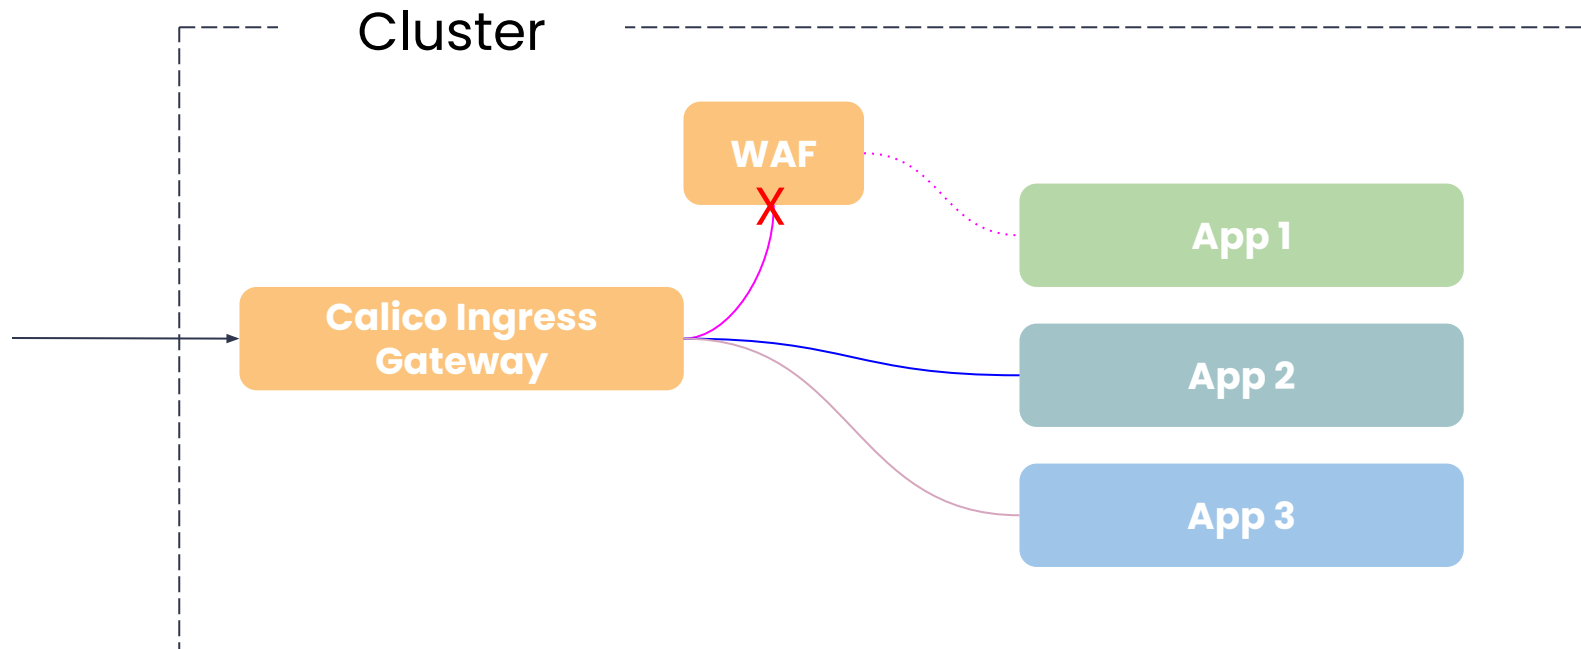
Circuit Breaking



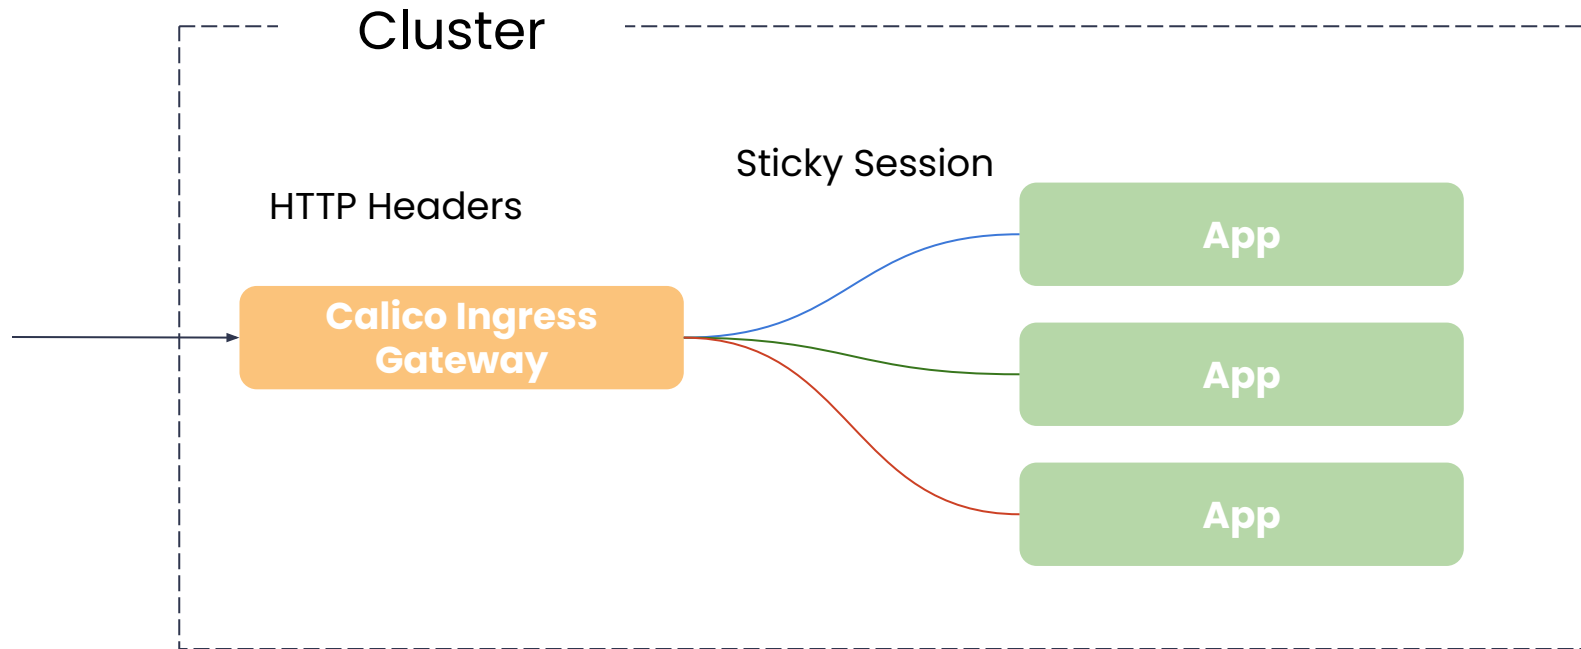
2.3

Other Traffic Functions

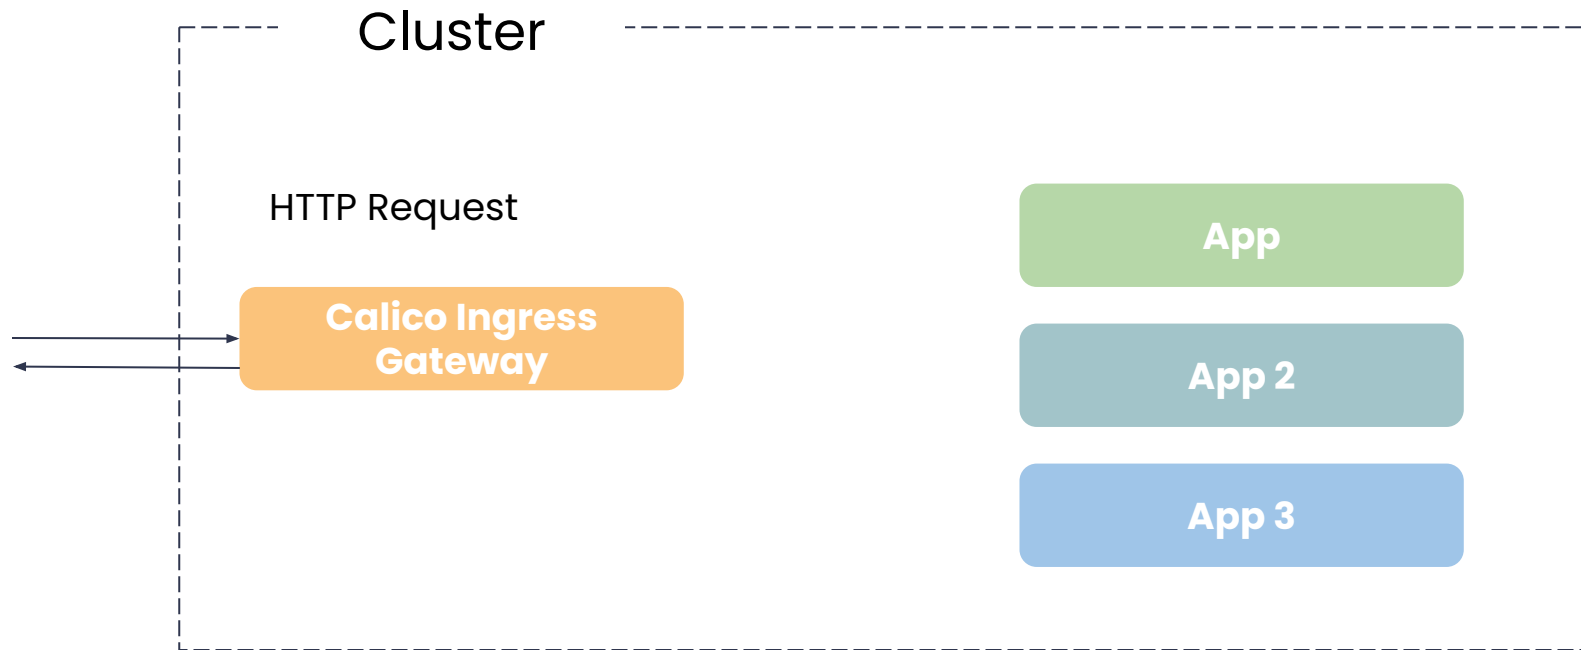
Coming Soon – WAF



Advanced Load Balancing – Consistent Hash



Direct Response



And lots and lots more...

- Circuit Breaker
- Backend Routing
- Client Traffic Policy
- Connection Limit
- Direct Response
- Failover
- Fault Injection
- GRPC Routing
- HTTP Redirects
- HTTP Request Headers
- HTTP Response Headers
- HTTP Timeouts
- HTTP URL Rewrite
- HTTP Request Mirroring
- Multicluster Service Routing
- Response Compression
- Response Override
- Retry
- Accelerated TLS Handshakes
- API Key Authentication
- Backend Mutual TLS: Gateway to Backend
- Backend TLS: Gateway to Backend
- Basic Authentication
- CORS
- External Authorization
- IP Allowlist/Denylist
- JWT Authentication
- JWT Claim-Based Authorization
- Mutual TLS: External Clients to the Gateway
- OIDC Authentication
- Threat Model
- TLS Passthrough
- TLS Termination for TCP

Thank you



TIGERA

Follow us on:

